

gradient descent

Kebing Liu

June 9

```
levels <- read.csv(file = "SMM_levels.csv", header = TRUE)

levels$log_tries_taken <- log(levels$tries_taken)
levels$Easy1 <- rep(0, length(levels$difficulty))
levels$Easy1[which(levels$difficulty == "Easy")] <- 1
```

Generate a contour plot

```
objective.fn <- function(alpha, beta) {
  # local variables
  x <- levels$log_tries_taken
  y <- levels$Easy1
  xbar <- mean(x)
  N <- length(x)

  # function on non-vector values
  non.vec <- function(a, b) {
    pi <- pnorm(a + b * (x - xbar))
    - 1/N * sum(y * log(pi / (1 - pi)) + log(1 - pi))
  }

  # apply non.vec
  n <- length(alpha)

  result <- rep(0, n)
  for (i in 1:n) {
    result[i] <- non.vec(alpha[i], beta[i])
  }
  return (result)
}


# generate alpha, beta
a <- seq(-1, 1.5, length.out = 50)
b <- seq(-2, 0, length.out = 50)

# cite from tutorial
# This is the key computation; there are many ways to achieve this.
z <- outer(a, b, objective.fn) # Outer takes advantage of the fact that
                               # objective.fn is vectorized simultaneously
                               # Matrix as heatmap.
image(a, b, z,
      col = heat.colors(100), # Palette with 100 levels; visually continuous.
      useRaster = TRUE,       # Less accurate image, but faster.
      ann = FALSE)
```

```

mtext(text = "alpha", side = 1, line = 1.5) # Suppress default axis labels, and
mtext(text = "beta", side = 2, line = 2) # draw them closer to the axis.
contour(a, b, z, add = TRUE) # Add contours

```



```

# modified from tutorial
gradientDescent <- function(theta = 0, rhoFn, gradientFn, lineSearchFn,
  testConvergenceFn, maxIterations = 100, tolerance = 1e-06, relative = FALSE,
  lambdaStepsize = 0, lambdaMax = 0.5) {
  # Pre-allocate a matrix to track theta at each iteration.
  SolutionPath = matrix(NA, nrow = maxIterations + 1, ncol = length(theta))
  SolutionPath[1, ] = theta
  converged <- FALSE
  i <- 0

  while (!converged & i <= maxIterations) {
    # use unnormalized gradient
    g <- gradientFn(theta)

    lambda <- lineSearchFn(theta, rhoFn, g, lambdaStepsize = lambdaStepsize,
      lambdaMax = lambdaMax)
    thetaNew <- theta - lambda * g
    converged <- testConvergenceFn(thetaNew, theta, tolerance = tolerance,
      relative = relative)
    theta <- thetaNew
  }
}

```

```

    # track the sequence of updates
    i <- i + 1
    SolutionPath[i + 1, ] = theta
  }
  ## Return last value and whether converged or not
  return(list(theta = theta, converged = converged, iteration = i, fnValue = rhoFn(theta),
    SolutionPath = SolutionPath[1:i, ]))
}

```

perform unnormalized gradient descent using the function

```

# use code from Q3
createObjProbit <- function(x, y) {
  ## local variable
  xbar <- mean(x)
  N <- length(x)

  ## return this function
  function(theta) {
    alpha <- theta[1]
    beta <- theta[2]
    pi <- pnorm(alpha + beta * (x - xbar))
    -1/N * sum(y * log(pi/(1 - pi)) + log(1 - pi))
  }
}

createGradientProbit <- function(x, y) {
  ## local variables
  xbar <- mean(x)
  N <- length(x)

  # return the function
  function(theta) {
    alpha <- theta[1]
    beta <- theta[2]
    yhat <- alpha + beta * (x - xbar)
    pi <- pnorm(yhat)

    # the gradient of -l(theta) is just negative gradient of
    # l(theta)
    -1/N * c(sum((y - pi)/(pi - pi^2) * dnorm(yhat) * 1), sum((y -
      pi)/(pi - pi^2) * dnorm(yhat) * (x - xbar)))
  }
}

### line searching could be done as a simple grid search
gridLineSearch <- function(theta, rhoFn, g, lambdaStepsize = 0.01, lambdaMax = 1) {
  ## grid of lambda values to search
  lambdas <- seq(from = 0, by = lambdaStepsize, to = lambdaMax)

  ## line search
  rhoVals <- sapply(lambdas, function(lambda) {
    rhoFn(theta - lambda * g)
  })
}

```

```

})
## Return the lambda that gave the minimum
lambdas[which.min(rhoVals)]
}

### Where testConvergence might be (relative or absolute)
testConvergence <- function(thetaNew, thetaOld, tolerance = 1e-10, relative = FALSE) {
  sum(abs(thetaNew - thetaOld)) < if (relative)
    tolerance * sum(abs(thetaOld)) else tolerance
}

# create rho and gradient functions
rho <- createObjProbit(levels$log_tries_take, levels$Easy1)
grad <- createGradientProbit(levels$log_tries_take, levels$Easy1)

# initialize starting values
starting.values = list(c(1.5, 0), c(-1, 0), c(1.5, -2))

# apply gradient search
Optim.list = lapply(starting.values, function(theta) {
  gradientDescent(rhoFn = rho, gradientFn = grad, theta = theta, lineSearchFn = gridLineSearch,
    testConvergenceFn = testConvergence, lambdaStepsize = 0.001, lambdaMax = 1,
    tolerance = 0.001)
})


# draw the contour plot
image(a, b, z, col = heat.colors(100), useRaster = TRUE)
contour(a, b, z, add = TRUE)

# draw the solution path
colour.list = c("purple", "gray", "cyan")

dummy = lapply(1:3, function(i) {
  solution.path = Optim.list[[i]]$SolutionPath
  n = nrow(solution.path)

  # draw the arrow
  for (j in 1:(n - 1)) {
    arrows(solution.path[j, 1], solution.path[j, 2], solution.path[j +
      1, 1], solution.path[j + 1, 2], length = 0.12, angle = 15,
      lwd = 3, col = adjustcolor(colour.list[i], alpha.f = 0.5))
  }
})
legend("bottomleft", legend = c("(1.5, 0)", "(-1, 0)", "(1.5, -2)"), col = colour.list,
  lwd = 3, text.col = "white")

```



```
mat = cbind(do.call(rbind, starting.values), do.call(rbind, lapply(Optim.list,
  unlist)))
knitr::kable(mat[, 1:7], booktabs = TRUE, col.names = c("\$\\alpha_0$",
  "\$\\beta_0$", "\$\\alpha_*$", "\$\\beta_*$", "Converged", "Iterations",
  "Obj. Value"))
```

α_0	β_0	α_*	β_*	Converged	Iterations	Obj. Value
1.5	0	0.3299263	-0.8581372	1	16	0.4158313
-1.0	0	0.3280023	-0.8568127	1	15	0.4158312
1.5	-2	0.3297111	-0.8579883	1	22	0.4158312

```
createStochasticGrad <- function(x, y, nsize = 10) {
  ## local variables
  N <- length(x)
  function(theta) {
    # generate samples
    subset = sample(N, nsize)
    x.new <- x[subset]
    y.new <- y[subset]

    alpha <- theta[1]
    beta <- theta[2]
    xbar <- mean(x.new)
```

```

    yhat <- alpha + beta * (x.new - xbar)
    pi <- pnorm(yhat)

    # the gradient of -l(theta) is just negative gradient of
    # l(theta)
    -1/N * c(sum((y.new - pi)/(pi - pi^2) * dnorm(yhat) * 1), sum((y.new -
    pi)/(pi - pi^2) * dnorm(yhat) * (x.new - xbar)))
  }
}

```

Perform random sample stochastic gradient descent

```

# the functions used for gradient descent use fixed step size
fixedLineSearch <- function(theta, rhoFn, g, lambdaStepsize = 0.01, lambdaMax = 1) {
  lambdaStepsize
}

gradientDescent <- function(theta = 0, rhoFn, gradientFn, lineSearchFn,
  testConvergenceFn, maxIterations = 100, tolerance = 1e-06, relative = FALSE,
  lambdaStepsize = 0, lambdaMax = 0.5) {
  # Pre-allocate a matrix to track theta at each iteration.
  SolutionPath = matrix(NA, nrow = maxIterations + 2, ncol = length(theta))
  SolutionPath[1, ] = theta
  converged <- FALSE
  i <- 0

  while (!converged & i <= maxIterations) {
    # use normalized gradient
    g <- gradientFn(theta)
    glength <- sqrt((g[1])^2 + (g[2])^2) ## gradient direction
    if (glength > 0)
      g <- g/glength

    lambda <- lineSearchFn(theta, rhoFn, g, lambdaStepsize = lambdaStepsize,
      lambdaMax = lambdaMax)
    thetaNew <- theta - lambda * g
    converged <- testConvergenceFn(thetaNew, theta, tolerance = tolerance,
      relative = relative)
    theta <- thetaNew

    # track the sequence of updates
    i <- i + 1
    SolutionPath[i + 1, ] = theta
  }

  ## Return last value and whether converged or not
  return(list(theta = theta, converged = converged, iteration = i, fnValue = rhoFn(theta),
    SolutionPath = SolutionPath[1:i, ]))
}

grad2 <- createStochasticGrad(levels$log_tries_taken, levels$Easy1, nsize = 25)
starting.values = list(c(1.5, 0), c(-1, 0), c(1.5, -2))

# perform stochastic gradient descent

```

```

Optim.list2 = lapply(starting.values, function(theta) {
  gradientDescent(rhoFn = rho, gradientFn = grad2, theta = theta, lineSearchFn = fixedLineSearch,
    testConvergenceFn = testConvergence, lambdaStepsize = 0.05, maxIterations = 500)
})

par(mfrow = c(1, 3))

# draw the contour
image(a, b, z, col = heat.colors(100))
contour(a, b, z, add = TRUE)

# draw the solution path
colour.list = c("purple", "gray", "cyan")

dummy = lapply(1:3, function(i) {
  solution.path = Optim.list2[[i]]$SolutionPath
  n = nrow(solution.path)

  # draw the arrow
  for (j in 1:(n - 1)) {
    arrows(solution.path[j, 1], solution.path[j, 2], solution.path[j +
      1, 1], solution.path[j + 1, 2], length = 0.12, angle = 15,
      lwd = 3, col = adjustcolor(colour.list[i], alpha.f = 0.5))
  }
  legend("bottomleft", legend = c("(1.5, 0)", "(-1, 0)", "(1.5, -2)"),
    col = colour.list, lwd = 3)
})

plot.iteration <- function(i, ...) {
  max <- max(c(Optim.list2[[1]]$SolutionPath[, i], Optim.list2[[2]]$SolutionPath[,
    i], Optim.list2[[3]]$SolutionPath[, i]))
  min <- min(c(Optim.list2[[1]]$SolutionPath[, i], Optim.list2[[2]]$SolutionPath[,
    i], Optim.list2[[3]]$SolutionPath[, i]))

  # plot the empty coordinate
  plot(1, type = "n", xlab = "iteration", xlim = c(0, 501), ylim = c(min,
    max), ...)

  color.list = c("purple", "gray", "cyan")

  for (j in 1:3) {
    lines(Optim.list2[[j]]$SolutionPath[, i], col = color.list[j])
  }

  # find the minimum fn_value and use the alpha beta
  min.index = 1
  if (Optim.list2[[1]]$fnValue > Optim.list2[[2]]$fnValue) {
    min.index = 2
  }
  if (min.index == 2 && Optim.list2[[2]]$fnValue > Optim.list2[[3]]$fnValue) {
    min.index = 3
  } else if (min.index == 1 && Optim.list2[[1]]$fnValue > Optim.list2[[3]]$fnValue) {
    min.index = 3
  }
}

```

```

abline(h = Optim.list2[[min.index]]$theta[i], col = "black")

legend("bottomleft", legend = c("(1.5, 0)", "(-1, 0)", "(1.5, -2)"),
      col = colour.list, lwd = 3)
}

# plot alpha vs. iteration
plot.iteration(1, ylab = bquote(alpha), main = "alpha vs. iterations")

# plot beta vs. iteration
plot.iteration(2, ylab = bquote(beta), main = "beta vs. iterations")

```

